
Interprocess Communication (IPC)



Methods

- ▶ Shared memory
- ▶ Shared files
- ▶ Signals
- ▶ Message queues
- ▶ Pipes
 - ▶ Ordinary pipes
 - ▶ Named pipes
- ▶ Sockets

Shared Memory

Shared Memory

- ▶ Unix kernel prohibits one process from accessing (reading, writing) memory belonging to another process
 - ▶ System V shared memory allows a process to read and/or write to memory created by another
 - ▶ POSIX shared memory alternative is more portable
- ▶ Advantages:
 - ▶ Random access: you can update a small piece in the middle of a data structure, rather than the entire structure
 - ▶ Efficiency: directly accessed as it resides in the user process memory
- ▶ Disadvantages:
 - ▶ No automatic synchronization as in pipes or message queues
 - ▶ Pointers are only valid within a given process. Thus, pointer offsets cannot be assumed to be valid across inter-process boundaries. This complicates the sharing of linked lists or binary trees.

Example

```
#define SHM_NAME "my-shm"

int sid = shm_open(SHM_NAME, O_CREAT | O_RDWR, 0666);

int size = 2 * sizeof(int);
ftruncate(sid, size);

if (fork() == 0) { // child process
    int* p = (int*)mmap(NULL, size, PROT_WRITE, MAP_SHARED, sid, 0);
    p[0] = 10;
    p[1] = 20;
} else { // parent process
    wait(NULL);

    int* p = (int*)mmap(NULL, size, PROT_READ, MAP_SHARED, sid, 0);
    int x = p[0];
    int y = p[1];
    printf("x + y = %d\n", x + y);

    shm_unlink(SHM_NAME);
}
```

Compile with `-lrt` flag:
`?> gcc shared-mem.c -o shared-mem -lrt`

Shared Files

Shared Files

- ▶ Files can be accessed by multiple processes
- ▶ Advantages:
 - ▶ Simple
 - ▶ Portable
- ▶ Disadvantages:
 - ▶ Files resident in external memory are generally
 - ▶ Data may be read by external users/programs if no measure for control of permission is taken

Example

```
#define FILENAME "/tmp/shared-file"

if (fork() == 0) { // child process
    int x = 10, y = 20;

    int fd = open(FILENAME, O_CREAT | O_WRONLY);
    write(fd, &x, sizeof(x));
    write(fd, &y, sizeof(y));
    close(fd);

} else { // parent process
    wait(NULL);

    int x, y;
    int fd = open(FILENAME, O_RDONLY);
    int n1 = read(fd, &x, sizeof(x));
    int n2 = read(fd, &y, sizeof(y));
    close(fd);

    printf("x + y = %d\n", x + y);

}
```

Signals

Signals

- ▶ A signal is a notification to a process indicating the occurrence of an event
- ▶ Signal is also called software interrupt and is not predictable to know its occurrence, hence it is also called an asynchronous event
- ▶ Signal can be specified with a number or a name
 - ▶ Some are well-defined: **SIGKILL** (9), **SIGALRM** (14), **SIGTERM** (15),...
 - ▶ Some are left for user to define: **SIGUSR1** (10), **SIGUSR2** (12)
- ▶ When a signal is sent
 - ▶ The OS interrupts the target process' normal flow of execution to deliver the signal
 - ▶ If the process has previously registered a signal handler, that routine is executed
 - ▶ **SIGKILL** and **SIGSTOP** cannot be caught or ignored

Usage

- ▶ Set the handler for a signal:
 - ▶ `signal_handler_t signal(int signum, signal_handler_t func)`
 - ▶ `func` can be:
 - ▶ A user-defined handler function
 - ▶ `SIG_IGN`: Ignore the signal
 - ▶ `SIG_DFL`: Use the default handler
 - ▶ Returns the previous signal handler function
- ▶ Send a signal:
 - ▶ `int kill(pid_t pid, int signum)`
- ▶ Suspend a process until a signal arrives:
 - ▶ `int pause()`

Examples

- ▶ Avoiding Ctrl+C and Ctrl+Z:

- ▶ `signal(SIGINT, SIG_IGN);`
`signal(SIGTSTP, SIG_IGN);`

- ▶ Synchronization of processes:

- ▶ `void sig_handler_parent(int signum) {`
 `printf("Parent: Received a response signal from child\n");`
`}`

- `void sig_handler_child(int signum) {`
 `printf("Child: Received a signal from parent\n");`
 `kill(getppid(), SIGUSR1);`
`}`

- `if (fork() == 0) { // child process`
 `signal(SIGUSR1, sig_handler_child);`
 `pause();`
`} else { // parent process`
 `signal(SIGUSR1, sig_handler_parent);`
 `sleep(1);`
 `kill(pid, SIGUSR1);`
 `pause();`
`}`

Message Queues

Message Queues

- ▶ A message queue is a linked list of messages stored within the kernel and identified by a message queue identifier
- ▶ New messages are added to the end of a queue
 - ▶ `int msgsnd(int msgid, const void* msg, size_t size, int flags)`
 - ▶ A message has a positive long integer type field, a non-negative length, and the actual data bytes (corresponding to the length)
- ▶ Messages can then be fetched by another process
 - ▶ `ssize_t msgrcv(int msgid, void* msg, size_t size, long type, int flags)`
 - ▶ Pass 0 as `type` to fetch the 1st message of any type, or a positive value for the desired message type

Example

```
#define MQ_PATHNAME      "/tmp/my-queue"
#define MQ_PROJECT_ID    18
#define MESSAGE_TYPE     5

struct message_t {
    long msg_type;
    int x, y;
};

key_t key = ftok(MQ_PATHNAME, MQ_PROJECT_ID);
int msgid = msgget(key, 0666 | IPC_CREAT);
if (msgid < 0) {
    // handle error...
}


if (fork() == 0) { // child process
    struct message_t msg = { MESSAGE_TYPE, 10, 20 };
    msgsnd(msgid, &msg, sizeof(msg), 0);
} else { // parent process
    struct message_t msg;
    msgrcv(msgid, &msg, sizeof(msg), MESSAGE_TYPE, 0);
    mq_close(msgid);

    printf("x + y = %d\n", msg.x + msg.y);
}
```

Pipes

Unix Unnamed Pipes

- ▶ Unidirectional, synchronous
- ▶ Allow parent-child processes to communicate with each other
- ▶ Creating a pipe:
 - ▶ `int pipe(int fd[2])`
 - ▶ Pass it an array of two integers
 - ▶ On success, zero is returned, `fd` contains two file descriptors
 - ▶ Use `read()`, `write()`, `close()` on these file descriptors
 - ▶ On error, `-1` is returned
 - ▶ `errno` is set appropriately
 - ▶ `perror()` produces a message on standard error output, describing last error encountered during a system call or library function



Example

```
int fd[2];
if (pipe(fd) < 0) {
    // handle error...
}

if (fork() == 0) { // child process
    int x = 10, y = 20;
    write(fd[1], &x, sizeof(x));
    write(fd[1], &y, sizeof(y));
    close(fd[1]);
} else { // parent process
    int x, y;
    read(fd[0], &x, sizeof(x));
    read(fd[0], &y, sizeof(y));
    close(fd[0]);

    printf("x + y = %d\n", x + y);
}
```

Named Pipes (aka FIFO)

- ▶ A named pipe is a first-in-first-out queue
 - ▶ Written by one process and read by another
 - ▶ Possible to be attached to multiple processes
- ▶ An extension to the traditional pipe concept on Unix which lasts only as long as the process
 - ▶ A named pipe can last beyond the life of the process
 - ▶ A named pipe appears as a special kind of file on the local storage which allows two or more processes to communicate with each other by reading/writing to/from it, in the same way as an ordinary file

Example

```
#define FIFO_PATH "/tmp/my-fifo"

if (mkfifo(FIFO_PATH, 0666) < 0) {
    // handle error...
}


if (fork() == 0) { // child process
    int x = 10, y = 20;
    int fd = open(FIFO_PATH, O_WRONLY);
    write(fd, &x, sizeof(x));
    write(fd, &y, sizeof(y));
    close(fd);
} else { // parent process
    int x, y;
    int fd = open(FIFO_PATH, O_RDONLY);
    read(fd, &x, sizeof(x));
    read(fd, &y, sizeof(y));
    close(fd);

    printf("x + y = %d\n", x + y);
}
```

Sockets

Sockets

- ▶ A socket is defined as an endpoint for communication on the Internet
 - ▶ Identified by IP address and port
 - ▶ Networking protocol routes the packet to the destined host
 - ▶ Different servers listens for requests on well known port
 - ▶ E.g., the socket 146.86.5.20:1625 refers to port 1625 on host 146.86.5.20
- ▶ Communication consists between a pair of sockets
 - ▶ TCP socket: reliable, connection oriented, stream based – more popular method
 - ▶ UDP socket: unreliable, connectionless, datagram based



Client/Server Model

- ▶ Form of communication used by all network applications
- ▶ One application initiates communication and the other accepts

Server

- ▶ Starts first
- ▶ Passively waits for contact from a client at a prearranged location
- ▶ Responds to requests, can handle multiple remote clients simultaneously



Client

- ▶ Starts second
- ▶ Actively contacts a server with a request
- ▶ Waits for response from server

The Socket Interface

▶ Berkeley Sockets API

- ▶ Originally developed as part of BSD Unix (under gov't grant)
 - ▶ BSD = Berkeley Software Distribution
- ▶ Now the most popular C/C++ API for writing applications over TCP/IP
 - ▶ Also emulated in other languages: Perl, Tcl/Tk, etc.
 - ▶ Also emulated on other operating systems: Windows, etc.

▶ Basic usage:


- ▶ TCP: a socket is like a file
 - ▶ servers (passive open) `listen()` and `accept()`
 - ▶ clients (active open) `connect()`
 - ▶ both sides can then `read()` and/or `write()` (or `send()` and `recv()`)
 - ▶ then each side must `close()`
- ▶ UDP: uses `sendto()` and `recvfrom()`



TCP

Server

1. Create transport endpoint for incoming connection request: `socket()`
2. Assign transport endpoint an address: `bind()`
3. Announce willing to accept connections: `listen()`
4. Block and wait for incoming request: `accept()`
5. Wait for a packet to arrive: `recv()`
6. Formulate reply (if any) and send: `send()`
7. Release transport endpoint: `close()`



Client

1. Create transport endpoint: `socket()`
2. Assign transport endpoint an address (optional): `bind()`
3. Determine address of server
4. Connect to server: `connect()`
5. Formulate message and send: `send()`
6. Wait for packet to arrive: `recv()`
7. Release transport endpoint: `close()`

UDP

Server

1. Create transport endpoint for incoming connection request: `socket()`
2. Assign transport endpoint an address: `bind()`
3. Wait for a packet to arrive: `recvfrom()`
4. Formulate reply (if any) and send: `sendto()`
5. Release transport endpoint: `close()`



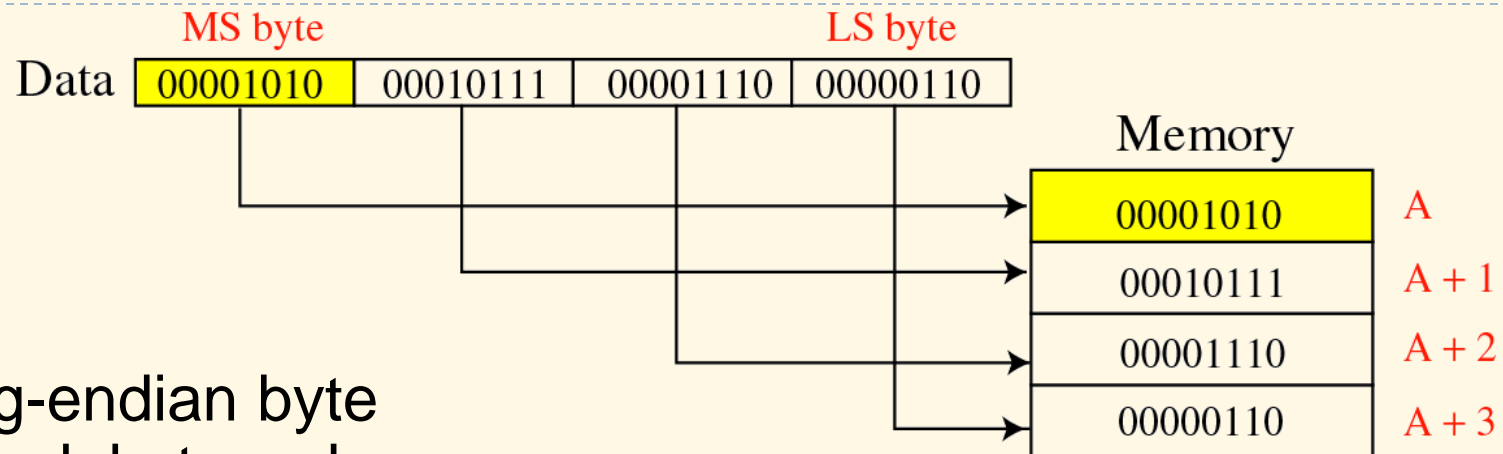
Client

1. Create transport endpoint: `socket()`
2. Assign transport endpoint an address (optional): `bind()`
3. Determine address of server
4. Formulate message and send: `sendto()`
5. Wait for packet to arrive: `recvfrom()`
6. Release transport endpoint: `close()`


Address Resolution

- ▶ `int getaddrinfo(
 const char* host,
 const char* service,
 const struct addrinfo* hints,
 struct addrinfo** res
);`
- ▶ Returns a linked list of one or more `addrinfo` structures, each of which contains an Internet address that can be specified in a call to `bind()` or `connect()`
- ▶ IPv4 and IPv6 compatible

Byte Ordering and Transformation



- ▶ TCP/IP uses big-endian byte order, aka network byte order
- ▶ Host \Leftrightarrow network byte order conversion is necessary
 - ▶ `uint16_t htons(uint16_t v)`
 - ▶ `uint16_t ntohs(uint16_t v)`
 - ▶ `uint32_t htonl(uint32_t v)`
 - ▶ `uint32_t ntohl(uint32_t v)`



Example

- ▶ See:
 - ▶ `tcp_server.c`
 - ▶ `tcp_client.c`